



CYBERSECURITY FOR PROFESSIONALS BOOTCAMP RUBRIC

Capstone Cybersecurity for Professionals Process

Purpose

The capstone serves as the culminating training activity where participants demonstrate their technical and professional skills through a team-based project and presentation before transitioning to on-the-job activities.

Capstone Planning

Introduce the Capstone Project

Timing: During trainer preparation and learner orientation. Or week before hybrid schedule.

Instructor Responsibilities

- Review the capstone project, deliverables, and rubric.
- Explain project expectations during the first week of training.
- Share examples of successful capstone presentations when available.
- Reinforce that the capstone is a major component of the final evaluation.

TAM / Program Manager Responsibilities

- Confirm the training team shares:
 - Capstone project documentation
 - Evaluation rubric
 - Presentation expectations
- Confirm anticipated capstone presentation dates with client stakeholders.
- Add capstone milestones to the program schedule.

Team Formation

Timing

Approximately 50-60% through the bootcamp. This allows instructors sufficient time to assess:

- Technical aptitude
- Participation
- Leadership potential
- Collaboration skills
- Attendance and professionalism

Team Assignment Guidelines

Recommended Team Size

Cohort Size	Team Structure
8-12 learners	Teams of 3-4
13-18 learners	Teams of 3-5

19-25 learners	Teams of 4-5
----------------	--------------

Team Formation Criteria

The instructor should create balanced teams based on:

Technical Mix

- Strong performers paired with developing learners
- Mix of technical strengths

Professional Skills Mix

- Strong communicators distributed across teams
- Emerging leaders spread across multiple teams

Engagement Levels

- Avoid placing all high performers on one team.
- Avoid placing all struggling learners together.

Role Assignment Each team should identify:

- Team Lead
- Technical Lead
- Presentation Lead
- Documentation Lead

One learner may hold multiple roles for smaller teams.

Ownership

Primary Owner: Instructor

The instructor has the best visibility into learner capabilities and should make final team assignments. TAM may provide input regarding behavioral or professionalism concerns

Capstone Execution

Timing

Final 1-2 weeks of training, depending on program length.

Instructor Responsibilities

- Manage day-to-day project execution.
- Conduct project check-ins.
- Coach teams through technical challenges.
- Monitor contribution levels across team members.
- Ensure all learners participate meaningfully.
- Review project artifacts and drafts.

TAM Responsibilities

- Monitor overall engagement.

- Address professionalism concerns.
- Support any learner escalations.
- Ensure project milestones remain on schedule.

Presentation Logistics

Primary Logistics Owner

Lead Instructor

The instructor manages:

- Presentation order
- Timing allocations
- Technical readiness checks
- Project evaluation
- Q&A facilitation

TAM Responsibilities

The TAM coordinates:

- Client invitations
- Observer attendance
- Calendar management
- Recording arrangements
- Distribution of presentation schedule

Current guidance already places presentation coordination and client logistics with the TAM while instructors facilitate the actual presentation sessions.

Presentation Standards

- Every learner must actively participate.
- Every learner must speak during the presentation.
- Recommended presentation slot:
 - 15-20 minute presentation
 - 5-10 minute Q&A
 - 5 minute transition between teams

For larger cohorts, multiple learners should present as a single project team rather than individual presentations. This approach is already referenced in trainer planning discussions.

Evaluation

Instructor Responsibilities

- Complete scoring rubric.
- Assess:
 - Technical quality
 - Problem solving

- Teamwork
- Presentation effectiveness
- Documentation quality

TAM Responsibilities

- Track outcomes.
- Identify readiness concerns.
- Document behavioral feedback.
- Share results with stakeholders as appropriate.

Evaluation Overview

This rubric is used to evaluate team capstone projects for the Cybersecurity for Professionals Bootcamp. Each team is scored across ten criteria spanning system hardening, network defense, cryptography, threat detection, incident response, data analytics, human-factor security, governance, documentation, and presentation. Each criterion is rated on a four-point scale and weighted to produce a final score out of 4.00.

Evaluators should anchor their scores to the descriptors in this rubric rather than to other teams. The descriptors are written so that a team can be Proficient in a category by delivering competent, professional work not by going beyond the assignment. Exemplary is reserved for work that an experienced SOC analyst or security engineer at a regulated bank would be comfortable standing behind in front of an examiner.

Scoring Scale

- 4 — Exemplary: meets the criterion at a level appropriate for an experienced security professional in a regulated banking environment.
- 3 — Proficient: meets the criterion at a level appropriate for an early-career security analyst or engineer.
- 2 — Developing: partially meets the criterion; gaps that would need to be closed before this work could be trusted in production.
- 1 — Beginning: does not meet the criterion; significant rework required.

A final score below 2.5 indicates the project does not meet the minimum bar for completion. A final score of 3.5 or above is presented to leadership as exemplary work.

Evaluation Criteria

The following table defines the descriptors for each criterion across the four-point scale.

Criterion	Exemplary (4)	Proficient (3)	Developing (2)	Beginning (1)
Network & System Hardening (Linux/Windows)	Hosts brought to and maintained in a known-good, least-privilege state.	Solid hardening across most hosts. Least privilege is mostly enforced. Minor gaps in	Some hardening steps completed, but reliance on defaults or manual	Hosts remain largely in default or vulnerable configuration. No coherent hardening approach.

Criterion	Exemplary (4)	Proficient (3)	Developing (2)	Beginning (1)
<i>Weight: 15%</i>	Hardening decisions are documented and clearly justified. Windows and Linux are treated with equal rigor.	documentation or consistency between platforms.	one-off fixes. Uneven coverage between Linux and Windows.	
Network Segmentation & Firewall Defense <i>Weight: 10%</i>	Segmentation and firewall rules clearly isolate the customer-facing tier from internal systems. Rules are minimal, justified, and mapped to an accurate diagram.	Reasonable segmentation and rule set. Diagram mostly matches implementation. Minor gaps in justification.	Some segmentation exists but boundaries are inconsistent or overly permissive. Diagram is incomplete or inaccurate.	Little or no meaningful segmentation. Firewall rules are default, absent, or not understood by the team.
Cryptography, PKI & Secure Communications <i>Weight: 10%</i>	Certificate correctly issued and configured with a fully explainable chain of trust. Team can speak fluently to the TLS handshake and key-management choices.	Certificate works and chain of trust is mostly understood. Small gaps in explanation of handshake or key management.	Certificate is issued but the team struggles to explain the trust chain or made avoidable configuration errors.	No working certificate, or the team cannot explain how TLS/PKI functions in their own environment.
Threat Detection & AI-Assisted SOC Triage <i>Weight: 15%</i>	Detection content reliably fires on the injected scenario. AI-assisted triage is used appropriately, with every AI-surfaced conclusion independently verified before action.	Detection content fires with minor gaps. AI triage is used and mostly verified, with occasional unverified assumptions.	Detection is partial — some activity is missed or only found after a hint. AI output is used with limited verification.	Detection content does not fire, or the team relies on being told about the incident rather than finding it. AI output is trusted without verification.
Incident Response, Fraud Investigation & Recovery <i>Weight: 15%</i>	Complete, accurate incident timeline from detection through contained, verified recovery. Fraud/BEC angle is correctly identified and addressed. Team narrates the response with confidence and technical depth.	Timeline is mostly complete and accurate. Fraud angle is identified. Recovery is verified with minor gaps in narration.	Timeline has gaps or sequencing errors. Fraud angle is only partially connected to the technical incident. Recovery is claimed but not clearly verified.	No coherent timeline. Fraud/BEC angle is missed entirely, or recovery is asserted without evidence.
Human Factor & Social Engineering Defense <i>Weight: 5%</i>	Social-engineering vector is clearly identified and addressed with both a technical control and a realistic awareness/process measure. Team explains the human element in plain, non-technical language.	Social-engineering vector is identified and addressed with at least one meaningful control. Explanation is mostly clear.	Social-engineering vector is mentioned but the remediation is generic or purely technical, missing the human element.	Social-engineering vector is not identified, or is dismissed without a proposed control.
Security & Fraud Analytics (Power BI Dashboard) <i>Weight: 10%</i>	Dashboard is built from real exercise data, clearly answers the operational question, and includes at least one well-used Copilot-generated visualization or summary. Governance guardrails (sensitivity labels, DLP, RLS) are applied and explained.	Dashboard is functional and readable, built from real data, with Copilot used at least once. Governance guardrails are mostly applied.	Dashboard exists but is difficult to read or uses placeholder/sample data not tied to the exercise. Limited or no Copilot use. Governance is mentioned but not applied.	No working dashboard, or the dashboard does not reflect the team's own exercise data.

Criterion	Exemplary (4)	Proficient (3)	Developing (2)	Beginning (1)
Governance, Risk & Compliance <i>Weight: 10%</i>	Compliance note clearly and correctly connects the incident response to a specific, relevant framework or guardrail (e.g., model risk management, data-privacy considerations). Team can defend the connection under questioning.	Compliance note references a relevant framework with a mostly correct connection to the incident.	Compliance note is present but generic, or the connection to the actual incident is weak.	No compliance note, or the framework referenced is irrelevant or misunderstood.
Documentation & Repository Quality <i>Weight: 5%</i>	README, architecture diagram, and incident record are clear, accurate, and current. A new analyst could understand what happened and how it was resolved from documentation alone.	Documentation is complete and reasonably clear, with minor inconsistencies.	Documentation covers the basics but has gaps or inaccuracies that would slow down a new analyst.	Documentation is missing, outdated, or misleading.
Presentation & Communication <i>Weight: 5%</i>	Presentation is well-paced, audience-aware, and appropriate for an organization-wide, non-technical-inclusive audience. Every team member speaks with technical depth. Q&A is handled confidently across the team.	Clear, well-organized presentation. Most decisions are explained well. Q&A is handled well by most of the team.	Presentation conveys the work but is uneven, or leans too technical/too shallow for the audience. Some Q&A reveals gaps in understanding.	Presentation is disorganized or shallow. Q&A surfaces significant knowledge gaps, or not every learner participates.

Scoring Summary

Use the table below to record the per-criterion score and compute the weighted total. The weighted score for each row is the raw score multiplied by the weight; the final score is the sum of the weighted scores divided by 100.

Criterion	Weight	Score (1–4)	Weighted Score
Network & System Hardening (Linux/Windows)	15%		
Network Segmentation & Firewall Defense	10%		
Cryptography, PKI & Secure Communications	10%		
Threat Detection & AI-Assisted SOC Triage	15%		
Incident Response, Fraud Investigation & Recovery	15%		
Human Factor & Social Engineering Defense	5%		
Security & Fraud Analytics (Power BI Dashboard)	10%		
Governance, Risk & Compliance	10%		
Documentation & Repository Quality	5%		

Criterion	Weight	Score (1–4)	Weighted Score
Presentation & Communication	5%		
Total	100%		Final Score (÷100)